



Latest developments in **pyop3**

Connor Ward

13 September 2023



What is `pyop3`?

A simple-ish example

What's the point?

What is pyop3?



- A programming language for mathematicians
- Comes with a compiler
- The language lets you express how to read and write from complicated data structures



- A **domain-specific** programming language for mathematicians **embedded in Python (like UFL)**
- Comes with a **just-in-time** compiler **that targets loopy and then C/CUDA/OpenCL**
- The language lets you express how to read and write from complicated data structures
- **Never need to create a PetscSection ever again!**




- FEM codes have diverse and complicated data structures
- These data structures also need to be accessed in non-trivial ways

A simple-ish example




1. Create data structures
2. Execute loop expressions acting on the data structures

Create a data layout for a 2 cell mesh



```
mesh_axis = Axis(  
    [  
        AxisComponent(2, "cells"),  
        AxisComponent(5, "edges"),  
        AxisComponent(4, "verts"),  
    ],  
    "mesh"  
    permutation=[3, 7, 0, 10, 5, 9, 2, 1, 6, 4, 8],  
)
```


Create a data layout for a 2 cell mesh



```
mesh_axis = Axis(  
    [  
        AxisComponent(2, "cells"),  
        AxisComponent(5, "edges"),  
        AxisComponent(4, "verts"),  
    ],  
    "mesh"  
    permutation=[3, 7, 0, 10, 5, 9, 2, 1, 6, 4, 8],  
)
```




Now make it a P3 function space




```
axes = (  
    AxisTree(mesh_axis)  
        .add_subaxis(Axis(1), mesh_axis.id, "cells")  
        .add_subaxis(Axis(2), mesh_axis.id, "edges")  
        .add_subaxis(Axis(1), mesh_axis.id, "verts")  
)
```


Now make it a P3 function space



```
axes = (  
    AxisTree(mesh_axis)  
        .add_subaxis(Axis(1), mesh_axis.id, "cells")  
        .add_subaxis(Axis(2), mesh_axis.id, "edges")  
        .add_subaxis(Axis(1), mesh_axis.id, "verts")  
)
```



Now let's do a loop





```
for every cell in the mesh:  
    collect DoFs found in the cell's closure  
    call a local kernel with these DoFs  
    scatter the result to a global vector
```



```
for every cell in the mesh:  
    collect DoFs found in the cell's closure  
    call a local kernel with these DoFs  
    scatter the result to a global vector
```

```
loop(  
    c := mesh.cells.index(),  
    kernel(func0[closure(c)], ...)  
)
```

Code generation!



```
void my_loop(double *func0, int *map0, int *map1, int *layout0, int *layout1, int *layout2, ...) {
    double t_0[10]; // to store the "packed" data

    for (int32_t i_0 = 0; i_0 < 2; ++i_0) { // loop over cells
        // pack cell DoFs
        t_0[0] = func0[layout0[i_0]];

        // pack edge DoFs
        for (int32_t i_5 = 0; i_5 < 3; ++i_5) { // loop over edges
            for (int32_t i_6 = 0; i_6 < 2; ++i_6) { // loop over edge DoFs
                j_3 = map0[i_0 * 3 + i_5]; // select the right edge
                t_0[i_5*2 + i_6 + 1] = func0[layout1[j_3] + i_6]; // pack DoF
            }
        }

        // pack vertex DoFs
        for (int32_t i_7 = 0; i_7 < 3; ++i_7) { // loop over vertices
            j_5 = map1[i_0 * 3 + i_7]; // select the right vertex
            t_0[i_7 + 7] = func0[layout2[j_5]]; // pack DoF
        }

        // execute the local kernel and unpack the result
        kernel(t_0, ...);
        ...
    }
}
```


What's the point?

But PyOP2 can already do this, why do we need pyop3?



Is it faster than PyOP2?

But PyOP2 can already do this, why do we need pyop3?



Is it faster than PyOP2?

No!

But PyOP2 can already do this, why do we need pyop3?



Is it faster than PyOP2?

No!

Is it as fast as PyOP2?

But PyOP2 can already do this, why do we need pyop3?



Is it faster than PyOP2?

No!

Is it as fast as PyOP2?

Not yet!

But PyOP2 can already do this, why do we need pyop3?



Is it faster than PyOP2?

No!

Is it as fast as PyOP2?

Not yet!

So why is it useful?

But PyOP2 can already do this, why do we need pyop3?



Is it faster than PyOP2?

No!

Is it as fast as PyOP2?

Not yet!

So why is it useful?

It's not.

But PyOP2 can already do this, why do we need pyop3?



Is it faster than PyOP2?

No!

Is it as fast as PyOP2?

Not yet!

So why is it useful?

~~It's not.~~

But PyOP2 can already do this, why do we need pyop3?



Is it faster than PyOP2?

No!

Is it as fast as PyOP2?

Not yet!

So why is it useful?

~~It's not.~~

Composability!



- Performance optimisation is not the main priority, expressibility is
- PyOP2 has limitations:
 - Single loop
 - Single kernel
 - No map composition
 - Extruded meshes require invasive code changes
- Can do more in fewer lines of code
 - **pyop3** compiler is ~ 1000 lines of code, PyOP2's is ~ 2500
 - No special casing for extruded



- Interior facet integrals:

```
loop(facet := mesh.interior_facets, kernel(func0[closure(support(facet))]))
```

- Multigrid:

```
closure(fine2coarse(fine_cell))
```





```
loop(v := mesh.vertices.index, [  
  loop(c := star(v).index, [  
    kernel1(dat1[closure(c)], "mat"),  
    kernel2(dat2[closure(c)], "vec")  
  ]),  
  solve("mat", "vec", dat3[v])  
])
```


- We could also try to rewrite SLATE
- Loop composition can enable certain tiling optimisations

And this composability will work with LOADS of data structures...

Example 1: extruded



Example 2: ragged





Useful for variable-layer extrusion and PIC.



- Like ragged (no picture sorry)
- Arbitrary sparsity is completely possible

Example 4: “swapping” axes





- **pyop3** is a DSL/compiler framework for writing kernels with non-trivial access patterns
- It can do everything PyOP2 can do, and more!
- WIP